



ETH AI CENTER



United
Nations

Practical Work - Project Report

SemUN: A Semantics-Powered Search Platform for the United Nations' Digital Library

Clément Sicard
csicard@ethz.ch
D-INFK, ETH Zurich

Supervised by:

Dr. Menna El-Assady¹, Dr. Sascha Langenbach¹, Catherine Pysden, MSc.²

¹ETH Zurich ²United Nations

August 28, 2023

Keywords:

Natural Language Processing (NLP), Named Entity Recognition (NER), Graph databases, Frontend, Network Visualization

1 Introduction

The United Nations Digital Library (UNDL) is a United Nations (UN) service that provides public access to a diverse range of UN documents: voting data, speeches, maps, and open access publications starting from 1979. All of these documents have been classified according to the UNBIS Thesaurus, a multilingual database of the controlled vocabulary used to describe UN documents and other materials in the Library's collection, with a more or less precise topic label.

The main idea of this project is to create an analysis platform, with a network visualization of documents from a subset of the documents in the UN Digital Library. The analytics platform includes a basic Named Entity Recognition (NER) system to extract mentioned entities from the documents. The visualization part will be a network visualization, leveraging the versatility of the structure of a graph to display the extracted insights. Both parts aim at improving the search of documents by implementing an analytics layer on top of the existing search engine from the digital library.

2 Motivation & Scope

In the short-term, the goal of this project is to provide an MVP of a potential future UN product, in close contact with UN staff to make it conform to their needs. Specifically, this project will focus – as a first iteration – on Catherine Pysden's suggestion, "Women in Peacekeeping".

The project's long-term scope is to be used as a search engine for UN staff, member-state delegates, and members of the public with an interest in UN topics. Potential future work could include extending the project to the whole UN Digital Library. It could, for instance, also suggest Thesaurus-compliant metadata for untagged documents to facilitate the work of Library staff.

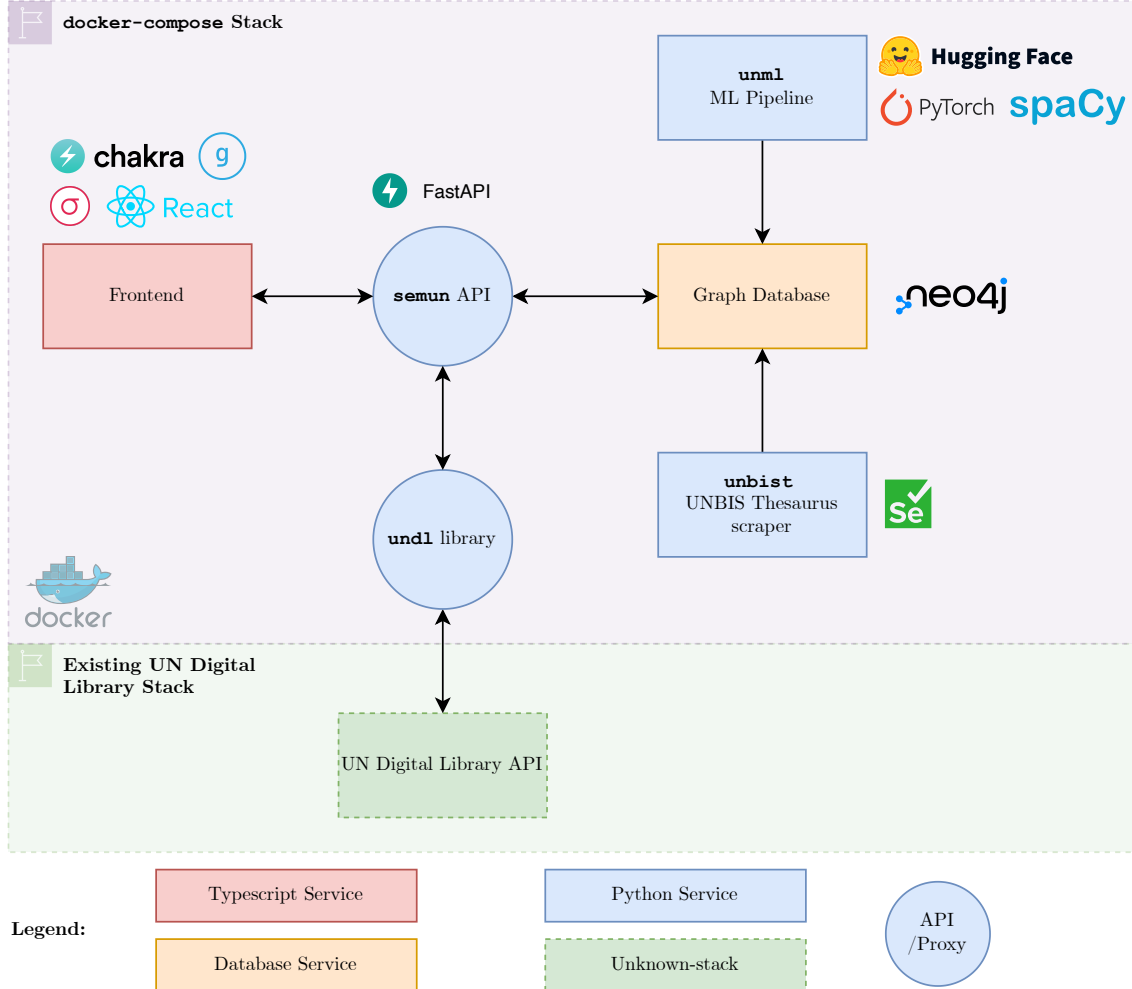





Figure 1: Final stack architecture

3 Final architecture

The final architecture is a full-stack architecture, from the database to the frontend, handed in as a Docker compose stack: **un-semun**.


The project consists of 7 GitHub repos:

- **un-semun**: The main repository with the **docker-compose** stack declaration.
- **un-semun-frontend**: The frontend.
- **un-semun-api**: The API for the frontend.
- **undl**: The code for **undl**, a Python wrapper around the UN Digital Library API.
- **un-unbis-thesaurus-scraper**: A scraper for the UNBIS Thesaurus taxonomy website.

-  **un-ml-pipeline**: Machine learning pipeline for UNDL documents.
-  **un-semun-misc**: Diverse scripts used for the project.
-  **un-semun-paper**: The code for this paper.

This paper will go through each of them in detail except for the paper repository.

3.1 un-semun-frontend: A React & Sigma.js frontend

I used React combined with Typescript for the UI framework, as well as Chakra UI for the UI components and **Sigma.js** via its React adapter **@react-sigma** for the network map. **graphology** was also used for graph manipulation in the frontend, mostly to iterate over graph elements to perform styling. The code is available here:  **un-semun-frontend**.

Add an actual screenshot of the frontend

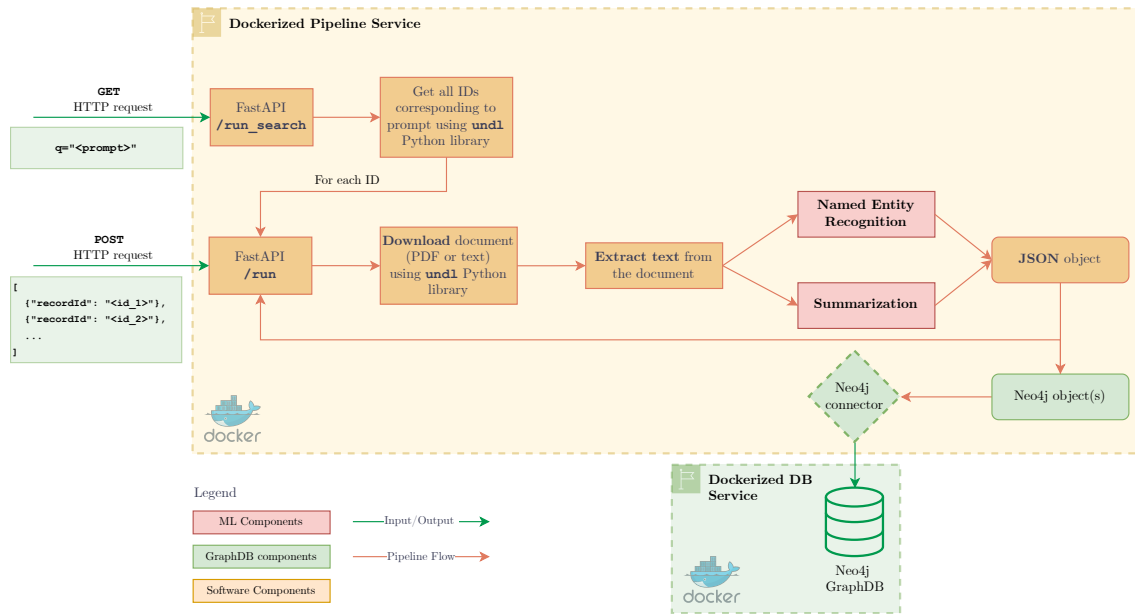


Figure 2: **un-semun-frontend** library: a dockerized machine learning pipeline for NER & summarization

The frontend was the part I was the least familiar with, but Chakra UI allowed to insert nice-looking components that I could customize based on my use case. It is composed in two panes:

- **The search bar** (on top): the user can enter its prompt, which will be sent to the API (3.2) to retrieve the results.
- **The result list** (on the left). The results are displayed as a scrollable list of **Card** components, with the title, the summary, and the date of publication. The user can click on a card to display the document in the right pane.
- **The network map** (on the right). The results of the search are also displayed as a network map, with the documents, related United Nations bodies, topics from UNBIS Thesaurus taxonomy, and named entities extracted from the documents. The user can click on a node to display the document in the left pane. This map is also fetched using the API (3.2) and is based on the results of the machine learning pipeline (3.5).

3.2 un-semun-api: An API for un-semun-frontend using FastAPI

The **un-semun-api** Python package was developed to provide an API for the frontend. It is a **FastAPI** application, which is a Python framework for building APIs. It is a modern framework, which is fast, easy to use, and well documented. One nice feature is that it is also coupled with **pydantic** to perform data validation and serialization. This is very useful to ensure that the data sent to the frontend is valid, and to avoid having to write boilerplate code for serialization and deserialization.

This package mainly acts as a proxy between the frontend and the UNDL API, and is composed of 2 main endpoints:


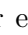
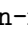

- **/search**: This endpoint is used to perform a search query. It takes a **GET** HTTP request with query parameter the prompt. Then, it uses **undl** (3.3) library to perform the search on the UNDL API and returns the results to the frontend.
- **/graph**: This endpoint is used to get the graph corresponding to the above search query, also with an HTTP **GET** request with the prompt as unique query parameter. To optimize the process, it first gets the IDs of all documents returned by the search query on the given prompt, then it queries the Neo4j graph database (3.6) to get the graph corresponding to these documents. Finally, it returns the graph to the frontend as **JSON** structured according to the format expected by the **graphology** Typescript library, which is used to model a graph and manipulate it as a programmatic object.

Note that the API is dockerized and that it offers a basic query caching mechanism to avoid querying the UNDL API too often. This is done using a simple **dict** in memory, which is not ideal but enough for the scope of this project.

3.3 undl: A Python library to wrap to the UN Digital Library API

For the **undl** library, I used Python 3.10 with packages, with **requests** for the HTTP requests, **pandas** for the data manipulation, and **pydantic** for the data validation.

Its purpose is to query the UN Digital Library API, and to return the results in a structured way and convenient way. It converts the **MARCXML** response from the API to a **JSON** object and implements caching to save the results of the queries, and to avoid querying the API again if the same query is made. The client offers these main methods:

- **query(...)**: Takes the prompt (e.g. "Women in peacekeeping") as an argument, and returns the detailed results of the query as a **JSON** object. The API response is a list of detailed documents (the fields we actually collect are precised in 3.6 for the **Document** node type). The main difference with a direct call to the UNDL API is that it converts the output to a **JSON** object. It is used by  **un-semun-api** (3.2).
- **getAllRecordIds(...)**: Same as **query**, but returns only the IDs of the documents, not their entire fields - hence is much faster. It is used by  **un-semun-api** (3.2) and  **un-ml-pipeline** (3.5).
- **queryBdId(...)**: Queries the API for a single document, given its **id**. It is used by  **un-ml-pipeline** (3.5).

These methods are callable either from the command-line interface (CLI), or from another Python script when imported as a library (see an example in Figure 3)

Note that **undl** requires a valid UNDL API key to work (set using the environment variable **UN_API**), because the API calls are authenticated by this 36 characters long key.



Figure 3: Querying the UNDL API from a Python script using `undl` library

3.4 `un-unbis-thesaurus-scraper`: the UNBIS Thesaurus scraper

🔗 `un-unbis-thesaurus-scraper` is a Python scraper that crawls the UNBIS Thesaurus website¹ and extracts the thesaurus terms and their relations. Once all the Thesaurus entries have been parsed, they are inserted with all their fields to the graph database instance (3.6).

It is mainly used to link Thesaurus topics the documents which are passed through the machine learning pipeline (3.5), but I felt it was also interesting to be able to visualize all the topics together and how they're related, so I built and deployed a small Vercel app to do so: ☁️ <https://un-graph-ui.vercel.app/>, which is built on a similar stack as the frontend (3.1).

3.5 `un-ml-pipeline`: The machine learning pipeline

🔗 `un-ml-pipeline` is the main component of this project. It englobes a Python library, `unml`, as well as an API wrapping it. It functions as follows (see Figure 4):

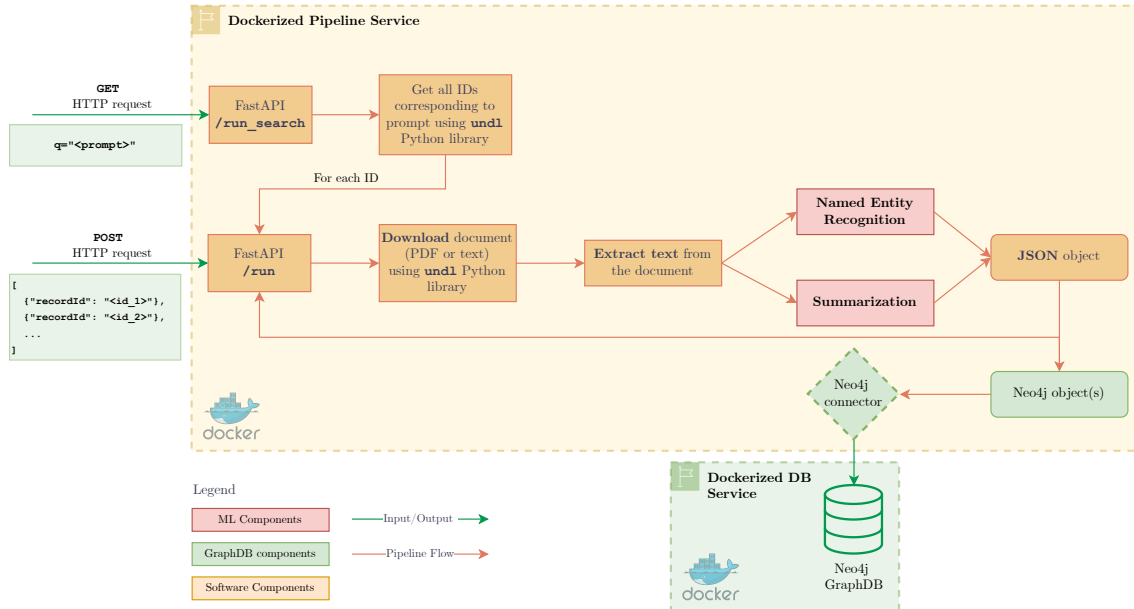


Figure 4: `unml` library: a dockerized machine learning pipeline for NER & summarization

The pipeline API is a `FastAPI` app, which exposes two endpoints. The `run` endpoint gets an ID as input, and works as follows:

¹<https://research.un.org/en/thesaurus>

1. A client sends an HTTP `POST` request to the pipeline API on `run` endpoint, with the payload containing a list of JSON objects with a `recordId`. (Note: this is arguably not an optimal data structure to store the IDs, but it was the easiest way to make `pydantic` happy and validate the input data structure.)
2. Use `undl`'s client `queryBdId` method to download the corresponding document complete information and PDF document. Sometimes it doesn't exist – if it is the case, jump to step ???
3. Extract the text from the PDF document using `PyMuPDF`² library.
4. Then the pipeline is split into 2 parts:
 - **Named Entity Recognition:** Extract named entities from the document. All types entities are collected (the types can vary depending on the model used though), but for data quality reasons, only UN bodies and Countries are actually stored in the graph database and linked to the document. However, the other entities are still returned by the API at the end of a pipeline run, and can be used for further analysis.
 - **Summarization:** The document is being summarized using a deep learning model, and the chosen model is a parameter of the pipeline (see 3.5.1 for more details). The summary for a document is then stored in the database to enhance the nodes.
5. The results of the machine learning components (e.g. the summary and the extracted named entities) are then stored into a JSON object
6. The JSON object is converted to a Cypher (the graph database, Neo4j, query language) query and then the query is run against the Neo4j database to update the graph. Note that the Neo4j instance lives in another Docker container to better compartmentalize the stack.

add step

The machine learning pipeline also offers the `run_search` endpoint, which receives a prompt as input (e.g. "Peacekeeping"), and works as follows:

1. A client sends an HTTP `GET` request to the pipeline API on `run_search` endpoint, with the prompt as a query parameter `q`.
2. The pipeline API queries UNDL API against the given prompt and collects document IDs corresponding to the search results using `getAllRecordIds` method from `undl`'s. It then returns the list of IDs.
3. For each record ID, do what querying `run` endpoint does

3.5.1 Models

For Named Entity Recognition tasks:











-

For Summarization:

-

Add challenges

²<https://github.com/pymupdf/PyMuPDF>

Model	File	Paper	Authors	Year	HuggingFace model
DistillBART-CNN	 DistillBARTCNN.py	arXiv 	Shleifer et al.	2020	Link
DistillBART-XSUM	 DistillBARTXSUM.py	arXiv 	Shleifer et al.	2020	Link
DistilPegasusCNN	 DistilPegasusCNN.py	arXiv 	Shleifer et al.	2020	Link
LongFormer	 LED.py	arXiv 	Beltagy et al.	2020	Link
LongT5	 LongT5.py	arXiv 	Guo et al.	2022	Link

3.5.2 API

API



Connection to Neo4j

3.6 Neo4j graph database

A graph database seemed to be the best choice to store the data extracted from the documents and the links between them, since it captures this relationship property and is able to retrieve it very efficiently. Neo4j seemed to be the most accessible one to use, but other solutions exist (such as ArangoDB or OrientDB).

Here is the chosen data model for this project:

3.6.1 Types of nodes

- **Document:** This node type represents a document in the UNDL. It contains the document's id from the library management system, its summary (created when the document was run through the machine learning pipeline), symbol (some internal UN classification for the document - for instance A/C.5/43/SR.50), publication date, title, URL, and publication location.
- **Topic:** This node type represents a topic in the sense of the UNBIS Thesaurus taxonomy. I made a scraper  `un-unbis-thesaurus-scraper` (3.4), and used it to retrieve all the topics from the taxonomy website, and inserted them into the database instance. Each topic has `name`, `cluster`, `id` fields, as well as label fields for each of the 6 official UN languages (`labelEn`, `labelFr`, `labelEs`, `labelAr`, `labelZh`, `labelRu`).
- **MetaTopic:** UNBIS Thesaurus is structured as a hierarchy, and a **MetaTopic** is one of the top-level topics in this hierarchy. It has the same fields as a **Topic** node, and there are 18 of them.³
- **Country:** This node type represents a country as a member state of the UN. It has the same fields as a **Topic** node, and there are 193 of them.⁴ The data to create these nodes was retrieved from the UN website, and the Cypher queries were generated using a Python script in  `un-semun-misc` (3.8).
- **UNBody:** This node type represents a body of the United Nations. A body of the UN is an organizational unit within the United Nations system, established to carry out specific functions ranging from peacekeeping and humanitarian aid to diplomatic negotiations and policy recommendations. Some famous examples include the Security Council, the World Health Organization (WHO), the International Monetary Fund

³The list of **MetaTopic** nodes from UNBIS Thesaurus website

⁴The list of the 193 UN member states from the UN website

(IMF) etc... As for the countries, the data was retrieved from the UN website, and the Cypher queries were generated using a Python script in `un-semun-misc` (3.8).

3.6.2 Types of relationships

These nodes are linked together by the following relationships:

- `-[REFERENCES]->`: This relationship type links a `Document` node to a `Country` or `Entity` node, and is self-explanatory: the document references the target entity. This relationship is extracted by the machine learning pipeline.
- `-[HAS_SUBTOPIC]->`: Links a `MetaTopic` node to a `Topic` node to explicit the hierarchy between the two nodes. It was created by the scraper.
- `-[IS_ABOUT]->`: Links a `Document` node to a `Topic` node to indicate that the `Document` has been classified by UN Digital Library staff as a document about the `Topic`. This relationship is extracted from the UNDL API response.
- `-[RELATED_TO]->`: Links a `Topic` node to another `Topic` node indicate a semantic link between the two topics. This relationship is extracted from the UNBIS Thesaurus website, using the scraper as well.

3.7 un-semun: The main repository

`un-semun` is a repository that englobes all other repositories as submodules. It also contains the `docker-compose` stack declaration, which points to Dockerfiles in submodules. They are updated using the `Makefile` when new commits are added to the submodules. The port forwarding definitions and environment variables are also declared here. This is the main entry point to run the whole stack.

3.8 un-semun-misc

This repository contains scripts for the both the UN member states and the UN bodies: it first extracts the data from an Excel file and a CSV file respectively, then cleans it, and finally inserts it into the Neo4j database after having written a Cypher query. The scripts are written in Python and directly use the `neo4j` Python connector library.

3.9 General tech stack notes

For all the Python components, the dependencies are managed using `poetry`. They are also all dockerized, and the whole stack is orchestrated using `docker-compose`.

4 Limitations

5 Discussion & future work

Need to run on all the documents, not just on the "Women in Peacekeeping" documents.

Models are kind of limited sometimes, would probably need to retrain them

Someone more experienced in frontend design could make it even better

Conclusion